




Dijkstra's Algorithm OpenMP and CUDA Parallelization

Presented by:
Jingyuan Wang and Jianbo Pei



Outline

- Serial Dijkstra's Algorithm Review
- OpenMP Parallelization
 - Approaches
 - Results
- CUDA Parallelization
 - Approaches
 - Results
- Conclusion - OpenMP vs CUDA
- Future Works



Serial Dijkstra's Algorithm

Dijkstra's Algorithm is an algorithm aiming at finding the shortest path between the source and other nodes in a graph.

Let V be a set of all the nodes, and S be a set of explored nodes.

– For each u in S we know the shortest path distance from s , $d(u)$.

■ Initially $S = \{s\}$, $d(s) = 0$; for all $v \neq s$, $d(v) = \infty$

■ While $S \neq V$


– **Select unexplored node v (in $V-S$) that minimizes “distance label”**

$$\pi(v) = \min_{e=(u,v): u \in S} (d(u) + \ell_e)$$

– Add v to S and set $d(v) = \pi(v)$

– **Update distance label $\pi(w)$ for all neighbors w of v**

■ EndWhile



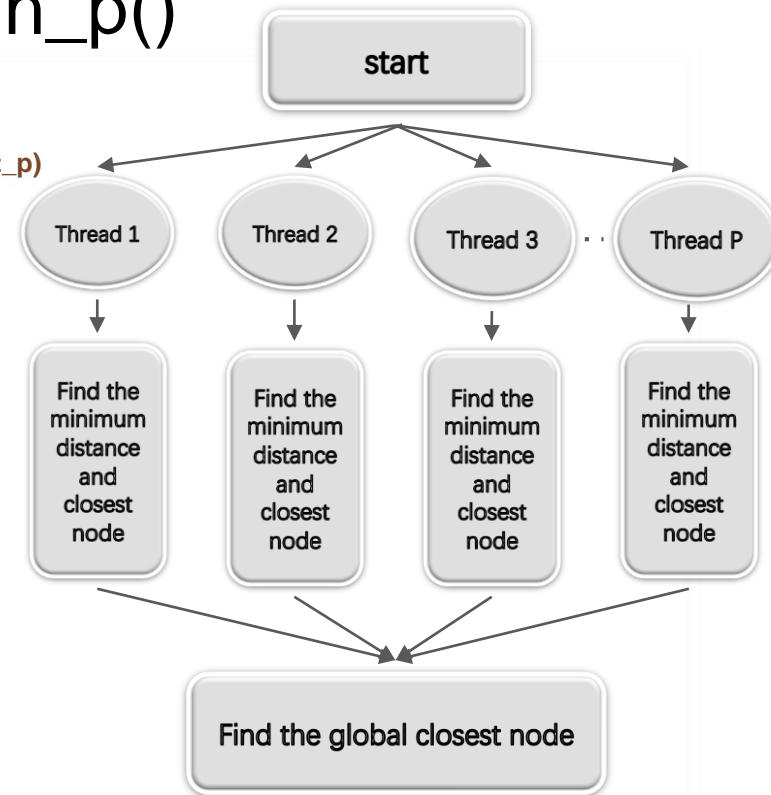
Two functions to parallel.



OpenMP Parallelization - findMin_p()

```
int findMin_p(int *dist_p, bool *visit_p){  
    omp_set_num_threads(THREADS);  
    #pragma omp parallel private(min_dist_thread, min_node_thread) shared(dist_p, visit_p)  
    {  
        min_dist_thread = min  
        min_node_thread = minNode  
        #pragma omp barrier  
        #pragma omp for nowait  
        for vertex = 0 to N with increment 4 {  
            if ((dist_p[vertex] < min_dist_thread) && (visit_p[vertex] == false))  
                Update min_dist_thread and min_node_thread  
            :  
        }  
        #pragma omp critical  
        {  
            if (min_dist_thread < min)  
                Update global closest node  
        }  
    }  
    return minNode;  
}
```

Unroll four iterations
of for loop.





OpenMP Parallelization - updateDist_p()

```
visit_p[minNode] = true;
void updateDist_p(){
    omp_set_num_threads(THREADS);
    #pragma omp parallel
    {
        #pragma omp for
        for vertex = 0 to N with increment 4 {
            if (the vertex has not been visited && the vertex connects to the minNode && distance[vertex] >
distance[minNode] + graph[minNode][vertex])
                distance[vertex] = distance[minNode] + graph[minNode][vertex];
            :
        }
        #pragma omp barrier
    }
}
```

Unroll four iterations
of for loop.




OpenMP - Results

Number of nodes	512	1024	2048	4096	8192	16384
Serial	0.005327	0.021351	0.084384	0.219032	0.867634	3.12067
<u>OpenMP</u>	<u>0.010836</u>	<u>0.022873</u>	<u>0.045601</u>	<u>0.111297</u>	<u>0.286892</u>	<u>0.745239</u>
Number of Threads	8	8	8	8	8	8

Highest improvement:
3.2x faster than serial



OpenMP - Results



CUDA Parallelization - Environment

- GPU: Tesla M2090 (512 cores, 6GB memory)
- Compile operations:
 - `module load cuda/5.0`
 - `module load gcc/4.4.3`
 - `nvcc -arch=sm_11 dijkstra_cuda.cu -o dijkstra_cuda`





CUDA Parallelization - closestNodeCUDA()

Step 1:

```
__global__ void closestNodeCUDA(int* min_value, int* minIndex, int* temp,

    unsigned int index = blockIdx.x * blockDim.x + threadIdx.x;

    __shared__ int cache[THREADS_PER_BLOCK];
    __shared__ int cacheIndex[THREADS_PER_BLOCK];

    if (index < num_vertices) {
        if ((node_dist[index]) < INT_MAX && (visited_node[index]) == 0) {
            cache[threadIdx.x] = node_dist[index];
            cacheIndex[threadIdx.x] = index;
        }
        else {
            cache[threadIdx.x] = INT_MAX;
            cacheIndex[threadIdx.x] = -1;
        }
    }
    __syncthreads();
```

__shared__:

cache[]: load distance values from global
node_dist array
cacheIndex[]: store corresponding indices



CUDA Parallelization - closestNodeCUDA()

Step 2:

```
unsigned int i = blockDim.x / 2;
while (i != 0) {
    if (threadIdx.x < i) {
        if (cache[threadIdx.x + i] < cache[threadIdx.x]) {
            cache[threadIdx.x] = cache[threadIdx.x + i];
            cacheIndex[threadIdx.x] = cacheIndex[threadIdx.x + i];
        }
    }
    __syncthreads();
    i /= 2;
}

if (threadIdx.x == 0) {
    temp[blockIdx.x] = cache[0];
    tempIndex[blockIdx.x] = cacheIndex[0];
}
```

reduction to find min and its index for each block.

store them at cache[0] and cacheIndex[0]

global memory:

temp[]: store all min from each block
tempIndex[]: store corresponding indices



CUDA Parallelization - closestNodeCUDA()

Step 3:

```
unsigned int k = BLOCKS / 2;
if (threadIdx.x == 0 && blockIdx.x == 0) {
    while (k != 0) {
        for (int j = 0; j < k; ++j) {
            if ((temp[j + k]) < temp[j]) {
                temp[j] = temp[j + k];
                tempIndex[j] = tempIndex[j + k];
            }
        }

        __syncthreads();
        k /= 2;
    }
}
```

reduction to find min and minIndex from global temp[] and global tempIndex[] using only thread 0 at block 0.

store min at temp[0]
store minIndex at tempIndex[0]



CUDA Parallelization - closestNodeCUDA()

Step 4:

```
if (threadIdx.x == 0 && blockIdx.x == 0) {  
    *min_value = temp[0];  
    *minIndex = tempIndex[0];  
  
    global_closest[0] = *minIndex;  
    visited_node[*minIndex] = 1;  
}  
__syncthreads();
```

finally, return min value from temp[0] and the corresponding index from tempIndex[0].

store minIndex to global_closest[0] which will be used in cudaRelax() function

mark this node as visited



CUDA Parallelization - cudaRelax()

```
__global__ void cudaRelax(data_t* graph, data_t* node_dist,  
    int next = blockIdx.x * blockDim.x + threadIdx.x;  
    int source = global_closest[0];  
  
    data_t edge = graph[source * VERTICES + next];  
    data_t new_dist = node_dist[source] + edge;  
  
    if ((edge != 0) &&  
        (visited_node[next] != 1) &&  
        (new_dist < node_dist[next])) {  
        node_dist[next] = new_dist;  
    }  
}
```

Update all nodes distance that are connected with global_closest node, which is returned from last function.




CUDA - Results

Number of nodes	512	1024	2048	4096	8192	16384
Serial	0.01	0.03	0.09	0.4	1.58	5.81
<u>CUDA</u>	<u>0.005334</u>	<u>0.010972</u>	<u>0.023383</u>	<u>0.051648</u>	<u>0.119978</u>	<u>0.330931</u>
Blocks	2	2	4	8	16	32
Threads/block	256	512	512	512	512	512

Highest improvement:
16.6x faster than serial



CUDA - Results







Conclusion - OpenMP vs CUDA



- In theory, running time complexity:
 - Serial: $O(N^2)$
 - OpenMP: $O(N * N/P)$
 - CUDA: $O(N * \log(N))$
- In our experiments, CUDA and OpenMP behave quite similar as above.





Future Works

- Looking for more optimization methods to improve current work.
- Testing on much larger graphs and real life graphs.





Thank you!



Questions?